Scheduling and Amortization for MPC

Benjamin Levy
levyb3@rpi.edu
Rensselaer Polytechnic Institute
Troy, New York

Benjamin Sherman
shermb@rpi.edu
Rensselaer Polytechnic Institute
Troy, New York

Lindsey Kennard
kennal@rpi.edu
Rensselaer Polytechnic Institute
Troy, New York

Ana L. Milanova
milanova@cs.rpi.edu
Rensselaer Polytechnic Institute
Troy, New York

Muhammad Ishaq[*]
m.ishaq@ed.ac.uk
University of Edinburgh
Edinburgh, Scotland

Vassilis Zikas[†]
vzikas@inf.ed.ac.uk
University of Edinburgh
Edinburgh, Scotland

ABSTRACT
CCS CONCEPTS
• **Theory of computation** → **Program analysis**; **Cryptographic protocols**; • **Security and privacy** → *Cryptography*.

KEYWORDS
protocol mixing; linear programming; multiparty computation; program analysis; cryptography

1 INTRODUCTION

2 BACKGROUND AND PROBLEM STATEMENT

2.1 Scheduling in MPC

Ana: Here we need more on MPC before jumping to scheduling. MPC-source, basic assumptions about static loop bounds, etc.

For this treatment we make the following simplifying assumptions:

(1) All statements in the program execute using the same protocol (sharing). That is, there is no share conversion. *Ishaq: This is not an assumption, this is our setting. We work in the single protocol world.*

(2) All MPC instructions have the same unit cost, 1 unit. *Ishaq: Don't we need to distinguish between local (cheap) vs. interactive (expensive) instructions? I don't see how we could assume this. – Past suggestion from Vassilis: Ideally we should assume cost to be a convex function. Ana: After discussion 1/12: replace single costs 1 with two levels of costs: α for non-local operations, e.g., MUL, and β for local ops, e.g., ADD.*

(3) There is unlimited bandwidth—i.e., a single MPC-instruction costs as much as N amortized instructions, namely 1 unit. *Ishaq: Comment from Vassilis: We assume there are infinite parallel capacities, this is the standard assumption in PRAM (Cryptographic Parallel RAM). Ana: PRAM assumption stays. We have to replace the unit cost of 1 with a suitable amortization function $f(n)$, which I don't think changes much in the cost modeling and analysis.*

(4) MPC instructions scheduled in parallel benefit from amortization *only if* they are the same instruction. Given our previous assumption, 2 MUL instructions scheduled in parallel benefit from amortization and cost 1, however a MUL and a MUX instructions scheduled in parallel still cost 2. *Ishaq: We need to reword this assumption to something like "K parallel MUL costs much less than K (because they get amortized), but any mix of K MUL and MUX still cost roughly K. Specifically, we should take away the low constants (2 and 1) because we know for low constants this is not true. Ana: Again, core assumption that MUL and MUX don't benefit from amortization stays. We have to replace constant costs with functions, as in (3).*

2.2 Scheduling in HPC

3 PRELIMINARIES

We assume arbitrarily nested loops in the MPC-source IR and read-only arrays. We assume that loops range from 0 to some constant N. Arrays are linearized (row-major order as in MOTION) and accesses are via functions of the induction variables of the enclosing loops. We write $i_1, i_2, \dots i_k$ to denote

the loop nest: i_1 is the outermost loop, i_2, is immediately nested in i_1, and so on until i_k. As an example, a statement nested in $i_1, i_2, ...i_k$ can access array $A[f(i_1, i_2, ...i_k)]$. We write $A[i_1, i_2, ...i_k]$ interchangeably. ANA: *Need to state this more precisely.*

3.1 Pseudo ϕ-nodes

A pseudo ϕ-node $X_1 = \phi(X_0, X_2)$ in a loop header is evaluated during circuit generation. If it is the 0-th iteration, then the ϕ-node evaluates to X_0, otherwise, it evaluates to X_2.

3.2 Def-use Edges

The dependence graph has the following def-use edges:

- same-level forward $X \to Y$ where X and Y are in the same loop nest $i_1, i_2, ...i_k$. E.g., $d = SUM(S[i, j], C[j])$ to $p = MUL(d, d)$ in Biometric is a same-level edge. A ϕ node can be a source of a same-level forward edge but not a target.
- outer-to-inner forward $X \to Y$ where X is in an outer loop nest, $i_1, i_2, ...i_j$, and Y is in an inner one, $i_1, i_2, ...i_j, ..., i_k$. A ϕ-node can be a source or a target of an outer-to-inner forward edge.
- inner-to-outer forward $X \to Y$ where X is a *phi*-node in an inner loop nest, $i_1, i_2, ...i_k, i_{k+1}$, and Y is in the enclosing loop nest $i_1, i_2, ...i_k$. E.g. $\text{sum}_0 = \phi(\text{sum}_1, 0)$ to $c = CMP(\text{sum}_0, \text{min}_0)$ is an inner-to-outer forward edge. Note that the source is *always* a ϕ-node in the immediately enclosing loop. The interpretation of this edge is that the use node Y uses the definition made in the last iteration of the inner loop. BEN: *The current representation of pseudo ϕ-nodes shows them attached to the loop header (i.e. in the loop). We may want to clarify that they are also evaluated at the loop's termination.*
- same-level back-edge $X \to Y$. Y is a phi-node in the header of the loop and X is a definition of the variable in the loop body. E.g., $\text{min}_1 = MUX(c, \text{sum}_1, \text{min}_1)$ to $\text{min}_0 = \phi(\text{min}_1, 10000)$ in Biometric is a same-level back-edge.
- inner-to-outer back-edge $X \to Y$: X and Y are both ϕ-nodes for some variable. The source X is in a loop nested into Y's loop (not necessarily immediately).
- mixed forward edge $X \to Y$. X is a *phi*-node in some loop $i_1, i_2, ...i_k, i_{k+1}$ and Y is a node in a loop nested into $i_1, i_2, ...i_k$. We transform mixed forward edges as follows. Let x_j be the variable defined at the phi-node X. We add a variable and assignment $\text{x}'_j = \text{x}_j$ immediately after the $i_1, i_2, ...i_k$ loop. Then we replace the use of x_j at Y with x'_j. This transforms a mixed forward edge into an "inner-to-outer" forward edge followed by an outer-to-inner forward edge. Thus Basic Vectorization handles one of "same-level", "inner-to-outer", or "outer-to-inner" def-use edges.

3.3 Helper Functions

We define $closure(n)$ where n is a phi-node. Intuitively, it computes the set of nodes (i.e., statements) that form a

dependence cycle with n. ANA: *Cycle(n) is probably a better name.* The closure of n is defined as follows:

- n is in $closure(n)$
- X is in $closure(n)$ if there is a same-level path from n to X, and $X \to n$ is a same-level back-edge.
- Y is in $closure(n)$ if there is a same-level path from n to Y and there is a same-level path from Y to some X in $closure(n)$.

We define the *raise_dim* (raise dimensions) and *drop_dim* (drop dimension) functions. Raise dimension "lifts" a lower-dimension array (ANA: *The right term here is tensor, I am pretty sure!*) into a higher dimension one. This is necessary when a lower-dimensional array is used in a higher dimensional loop and, essentially, is just copying of values. For example, Biometric contains the statement $d = SUB(S[i, j], C[j])$, in the j-loop which is nested into the i-loop. Here C is a one-dimensional array, however, to vectorize across both loops it is necessary to turn it into a two-dimensional array: $C[i, j]$ becomes $[C[0], C[1], ...C[J], C[0], C[1], ...C[J], ...C[0], C[1], ...C[J]]$, which turns the row into a matrix of I identical rows. (We use capital letters to denote the upper bounds of loops, e.g., J is the upper bound of the j-loop and I is the upper bound of the i-loop.)

$raise\_dim(A[i_1, ..., i_k], i_j)$ is defined as follows. It results in a new k+1-dimensional array A' where for every $0 \le i_k < I_{k+1}$, $A'[i_1, ..., i_j, ..., i_k] = A[i_1, ..., i_k]$. Adding n dimensions is trivially extended as a composition of n *raise_dim* that each adds a single dimension.

As expected, drop dimension turns a higher-dimensional array into a lower-dimensional one. The key use case is an inner-to-outer def-use edge. The code may define a variable, e.g., x in an inner loop, say j, then use this variable in an enclosing loop, say i. Our algorithm may vectorize the computation of x in the j-loop thus producing a vector $\text{x}[j]$ where $\text{x}[1]$ is the value of x after the 1st iteration and so on. Drop dimension states that the outer loop will use the value of the variable at the last iteration.

$drop\_dim(A[i_1, ..., i_j, ..., i_k], i_j)$ produces a k-dimensional array A' where $A'[i_1, ..., i_k] = A'[i_1, ..., I_j - 1, ..., i_k]$. In our analysis, dimensions are always dropped at the end, and again, one can define dropping n dimensions as a composition of n *drop_dim*.

BEN: *I'm not sure if I wrote up the above paragraphs correctly – the idea is that we're just copying values to extend over dimension j or only retain its last elements. The current writeup implies that $1 < j < k$ though, which isn't always true. Python implementations of raise_dim() and drop_dim() can be found here: https://github.com/milana2/ParallelizationForMPC/blob/master/compiler/compiler/motion_backend/reference_implementations.py*

ANA: *This section is still informal. Needs work to make more precise.*

4 VECTORIZATION

4.1 Basic Vectorization

{ Phase 1: Raise dimension of scalar variables to corresponding loop nest. We can traverse stmts linearly in MPC-source. }

for each def-use edge $stmt'(\text{def of } Y) \rightarrow stmt(\text{def of } X)$ **do**

    case def-use edge $stmt'(\text{def of } Y) \rightarrow stmt(\text{def of } X)$ of

        same-level: Y' is Y

        outer-to-inner: add $Y'[i, j, k] = raise\_dim(Y)$ at $stmt'$

        inner-to-outer: add $Y'[i, j, k] = drop\_dim(Y)$ at $stmt$

end for

{ Optimistically vectorize all. $\vec{i}$ means vectorized dimension. }

for each MPC $stmt : X = Op(Y_1, Y_2)$ in loop i, j, k **do**

    change to $X[\vec{i}, \vec{j}, \vec{k}] = Op(Y_1'[\vec{i}, \vec{j}, \vec{k}], Y_2'[\vec{i}, \vec{j}, \vec{k}])$

end for

{ Phase 2: Recreating FOR loops for cycles; vectorizable statements hoisted up. }

for each dimension d from highest to 0 **do**

    for each ϕ-node n in loop $i_1, ..., i_d$ **do**

        compute $closure(n)$

    end for

    {cl_1 and cl_2 intersect if they have common statement or update same array; "intersect" definition can be expanded}

    while there are closure cl_1 and cl_2 that intersect **do**

        merge cl_1 and cl_2

    end while

    for each closure cl (after merge) **do**

        create FOR $i_d = 0; ...$ loop

        add ϕ-nodes in cl to header block

        add target-less ϕ-node for A if cl updates array A

        add statements in cl to loop body in some order of dependences

        { Dimension is not vectorizable: }

        change $\vec{i_d}$ to i_d in all statements in loop

        treat FOR loop as monolith node: some def-use edges become same-level.

    end for

    for each target-less ϕ-node $A_1 = \phi(A_0, A_k)$ **do**

        in vectorizable stmts, replace use of A_1 with A_0

        discard ϕ-node if not used in any cl

    end for

end for

{ Phase 3: Dimensionality reduction: removing unnecessary dimensionality. }

{ A dimension i is dead on exit from stmt $X[...i...] = ...$ if all def-uses with targets outside of the enclosing FOR $i = 0...$ MOTION loop end at target (use) $X' = drop\_dim(X, i)$. }

for each stmt and dimension $X[...i...] = ...$ **do**

    if i is a dead dimension on exit from stmt $X[...i...] = ...$, remove i from X (all defs and uses)

end for

{ Now clean up drop_dim and raise_dim }

for each $X' = drop\_dim(X, i)$ **do**

    replace with $X' = X$ if i is dead in X.

end for

do (1) (extended) constant propagation, (2) copy propagation and (3) dead code elimination to get rid of redundant variables and raise and drop dimension statements

{ Phase 4: }

add SIMD for simdified dimensions

4.2 Example: biometric

We start from Benjamin's code with linear loops (MPC Source):

```
min_sum!1 = 10000
min_index!1 = 0
for i in range(0, N!0):
   min_sum!2 = PHI(min_sum!1, min_sum!4)
   min_index!2 = PHI(min_index!1, min_index!4)
   sum!2 = 0
   for j in range(0, D!0):
     sum!3 = PHI(sum!2, sum!4)
     d!3 = (S!0[((i * D!0) + j)] - C!0[j])
     p!3 = (d!3 * d!3)
     sum!4 = (sum!3 + p!3)
   !1!2 = (sum!3 < min_sum!2)
   min_sum!3 = sum!3
   min_index!3 = i
   min_sum!4 = MUX(!1!2, min_sum!3, min_sum!2)
   min_index!4 = MUX(!1!2, min_index!3, min_index!2)
!2!1 = (min_sum!2, min_index!2)
```

4.2.1 Phase 1 of Basic Vectorization

The transformation preserves the dependence edges. It raises the dimensions of scalars and optimistically vectorizes all operations. The next phase discovers loop-carried dependences and removes affected vectorization.

In the code below, all initializations (e.g., `min_sum!3 = i`), operations, and PHI nodes are *implicitly vectorized*. *raise_dim* and *drop_dim* statements, as well as propagation statements

(e.g., `min_sum!3 = sum!3^`) may have slightly different inter-pretation. ANA: *This is vague, may need elaboration.*

Note the two different versions of *raise_dim*. ANA: *More here! One just adds a least-significant dimension. The other one may reshape an input array.*

```
min_sum!1 = 10000
min_sum!1^ = raise_dim(min_sum!1, (i:N!0))
min_index!1 = 0
min_index!1^ = raise_dim(min_index!1, (i:N!0))
S!0^ = raise_dim(S!0, ((i * D!0) + j), (i:N!0,j:D!0))
C!0^ = raise_dim(C!0, j, (i:N!0,j:D!0))
for i in range(0, N!0):
   min_sum!2 = PHI(min_sum!1^, min_sum!4)
   min_index!2 = PHI(min_index!1^, min_index!4)
   sum!2 = 0 // Will lift, when hoisted
   sum!2^ = raise_dim(sum!2, (j:D!0)) // Special form?
   for j in range(0, D!0):
      sum!3 = PHI(sum!2^, sum!4)
      d!3 = S!0^ - C!0^
      p!3 = (d!3 * d!3)
      sum!4 = (sum!3 + p!3)
   sum!3^ = drop_dim(sum!3)
   !1!2 = (sum!3^ < min_sum!2)
   min_sum!3 = sum!3^
   min_index!3 = i  // Same-level, will lift when hoisted
   min_sum!4 = MUX(!1!2, min_sum!3, min_sum!2)
   min_index!4 = MUX(!1!2, min_index!3, min_index!2)
min_sum!2^ = drop_dim(min_sum!2)
min_index!2^ = drop_dim(min_index!2)
!2!1 = (min_sum!2^, min_index!2^)
```

4.2.2 Phase 2 of Basic Vectorization

This phase analyzes statements from the innermost loop to the outermost. The key point is to discover loop-carried dependencies and re-introduce loops whenever dependencies make this necessary.

Starting at the inner phi-node `sum!3 = PHI(sum!2, sum!4)`, the algorithm first computes its closure. The closure amounts to the phi-node itself and the addition node `sum!4 = (sum!3 + p!3)`, accounting for the loop-carried dependency of the computation of `sum`. The algorithm replaces this closure with a FOR loop on j removing vectoriaztion on j. Note that the SUB and MUL computations remain outside of the loop as they do not depend on phi-nodes that are part of cycles. The dependences are from `p!3[I,J] = (d!3[I,J] * d!3[I,J])` to the monolithic FOR loop and from the FOR loop to `sum!3 ≜ drop_dim(sum!3)`. (Lower case index, e.g., `i`, indicates non-vectorized dimension, while uppercase index, e.g., `I` indicates vectorized dimension.)

After processing inner loop code becomes:

```
min_sum!1 = 10000
min_sum!1^ = raise_dim(min_sum!1, (i:N!0))
min_index!1 = 0
min_index!1^ = raise_dim(min_index!1, (i:N!0))
S!0^ = raise_dim(S!0, ((i * D!0) + j), (i:N!0,j:D!0))
C!0^ = raise_dim(C!0, j, (i:N!0,j:D!0))
for i in range(0, N!0):
  min_sum!2[I] = PHI(min_sum!1^[I], min_sum!4[I])
  min_index!2[I] = PHI(min_index!1^[I], min_index!4[I])
  sum!2 = [0,..,0]
```

```
  sum!2^ = raise_dim(sum!2, (j:D!0))
  d!3[I,J] = S!0^[I,J] - C!0^[I,J]
  p!3[I,J] = (d!3[I,J] * d!3[I,J])
  for j in range(0, D!0):
    sum!3[I,j] = PHI(sum!2^[I,j], sum!4[I,j-1])
    sum!4[I,j] = (sum!3[I,j] + p!3[I,j])
  sum!3^ = drop_dim(sum!3)
  !1!2[I] = (sum!3^[I] < min_sum!2[I])
  min_sum!3 = sum!3^
  min_index!3 = i
  min_sum!4[I] = MUX(!1!2[I], min_sum!3[I], min_sum!2[I])
  min_index!4[I] = MUX(!1!2[I], min_index!3[I], min_index!2[I])
min_sum!2^ = drop_dim(min_sum!2)
min_index!2^ = drop_dim(min_index!2)
!2!1 = (min_sum!2^, min_index!2^)
```

When processing the outer loop two closures arise, one for `min_sum!2[I] = PHI(...)` and one for `min_index!2[I] = PHI(...)`. Since the two closures *do not* intersect, we have two distinct FOR-loops on i:

```
min_sum!1 = 10000
min_sum!1^ = raise_dim(min_sum!1, (i:N!0))
min_index!1 = 0
min_index!1^ = raise_dim(min_index!1, (i:N!0))
S!0^ = raise_dim(S!0, ((i * D!0) + j), (i:N!0,j:D!0))
C!0^ = raise_dim(C!0, j, (i:N!0,j:D!0))

sum!2 = [0,..,0]
sum!2^ = raise_dim(sum!2, (j:D!0))
d!3[I,J] = S!0^[I,J] - C!0^[I,J]
p!3[I,J] = (d!3[I,J] * d!3[I,J])

for j in range(0, D!0):
  sum!3[I,j] = PHI(sum!2^[I,j], sum!4[I,j-1])
  sum!4[I,j] = (sum!3[I,j] + p!3[I,j])

sum!3^ = drop_dim(sum!3)
min_index!3 = [0,1,2,...N!0-1] // or min_index!3 = [i, (i:N!0)]
min_sum!3 = sum!3^

for i in range(0, N!0):
  min_sum!2[i] = PHI(min_sum!1^[i], min_sum!4[i-1])
  !1!2[i] = (sum!3^[i] < min_sum!2[i])
  min_sum!4[i] = MUX(!1!2[i], min_sum!3[i], min_sum!2[i])

for i in range(0, N!0):
  min_index!2[i] = PHI(min_index!1^[i], min_index!4[i-1])
  min_index!4[i] = MUX(!1!2[i], min_index!3[i], min_index!2[i])

min_sum!2^ = drop_dim(min_sum!2)
min_index!2^ = drop_dim(min_index!2)
!2!1 = (min_sum!2^, min_index!2^)
```

4.2.3 Phase 3 of Basic Vectorization

This phase removes redundant dimensionality. It starts by removing redundant dimensions in MOTION loops followed by removal of redundant drop dimension statements. It then does (extended) constant propagation to "bypass" raise state-ments, copy propagation and dead code elimination.

The code becomes closer to what we started with:

```
min_sum!1 = 10000
min_index!1 = 0
```

```
S!0^ = raise_dim(S!0, ((i * D!0) + j), (i:N!0,j:D!0))
C!0^ = raise_dim(C!0, j, (i:N!0,j:D!0))

sum!2 = [0,..,0]
d!3[I,J] = S!0^[I,J] - C!0^[I,J]
p!3[I,J] = (d!3[I,J] * d!3[I,J])

// j is redundant for sum!3 and sum!4
for j in range(0, D!0):
  sum!3[I] = PHI(sum!2^[I], sum!4[I])
  sum!4[I] = (sum!3[I] + p!3[I,j])

// drop_dim is redundant, removing
// then copy propagation and dead code elimination
min_index!3 = [0,1,2,...N!0-1] // or min_index!3 = [i, (i:N!0)]

// i is redundant for min_sum!2, min_sum!4 but not for !12![i]!
for i in range(0, N!0):
  min_sum!2 = PHI(min_sum!1, min_sum!4)
  !1!2[i] = (sum!3[i] < min_sum!2)
  min_sum!4 = MUX(!1!2[i], sum!3[i], min_sum!2)

// same, i is redundant for min_index!2, min_index!4
for i in range(0, N!0):
  min_index!2 = PHI(min_index!1, min_index!4)
  min_index!4 = MUX(!1!2[i], min_index!3[i], min_index!2)

// drop_dim becomes redundant
!2!1 = (min_sum!2, min_index!2)
```

4.2.4 Phase 4 of Basic Vectorization

And this phase adds SIMD operations:

```
min_sum!1 = 10000
min_index!1 = 0
S!0^ = raise_dim(S!0, ((i * D!0) + j), (i:N!0,j:D!0))
C!0^ = raise_dim(C!0, j, (i:N!0,j:D!0))

sum!2 = [0,..,0]
d!3[I,J] = SUB_SIMD(S!0^[I,J],C!0^[I,J])
p!3[I,J] = MUL_SIMD(d!3[I,J] * d!3[I,J])

for j in range(0, D!0):
  // I dim is a noop. sum is already a one-dimensional vector
  sum!3[I] = PHI(sum!2^[I], sum!4[I])
  sum!4[I] = ADD_SIMD(sum!3[I],p!3[I,j])

min_index!3 = [0,1,...N!0-1]

for i in range(0, N!0):
  min_sum!2 = PHI(min_sum!1, min_sum!4)
  !1!2[i] = CMP(sum!3[i],min_sum!2)
  min_sum!4 = MUX(!1!2[i], sum!3[i], min_sum!2)

for i in range(0, N!0):
  min_index!2 = PHI(min_index!1, min_index!4)
  min_index!4 = MUX(!1!2[i], min_index!3[i], min_index!2)

!2!1 = (min_sum!2, min_index!2)
```

ANA: TODO: It will be good to add a pass in Phase 3, or post Phase 3, to get rid of extra dimensionality. Once we have vectorized as much as possible, there is no need to keep higher dimensionality. E.g., there is no need for the 2-dimensional array sum in the first for loop, and similarly, there is no need for 1-dimensional arrays for min_sum and min_index. We should get rid of these before generating MOTION code.

ANA: I don't have a crisp algorithm in mind. I don't think it will be difficult but we have to think this through. Put on our TODO list for us to think about since it will simplify MOTION code tremendously.

ANA: Begin Work-in-progress notes on implementation. Just some thoughts...

For now, I suggest we go with a "naive" implementation and try to do only minimal optimizations. All arrays are stored in the program as one-dimensional arrays. Arrays lifted from scalars have dimensionality of the loop enclosure in the original code, e.g., sum!3 is two dimensional [N!0,D!0] (but represented linearly in row-major order; all these "copies" are arrays of references to shares). Naturally, all writes into these arrays are canonical writes, i.e., $A[i, j] = ...$ and no things like $A[i+1, j-1] = ...$

The loop "for j in range..." is perhaps the most interesting. First, consider "projection", e.g., p!3[I,j]. We fix j (the current iteration) and iterate over I using the row-major formula; it returns the j-th column, a one-dimensional column vector. I think this can be generalized across any mix of vectorized vs. non-vectorized dimensions but I'm not 100% sure.

```
// project the (j-1)-st column vector:
A = sum!4[I,j-1]
// project j-th column from p!3:
B = p!3[I,j]
// we get a N!0-size vector of references to NEW shares:
C = ADD_SIMD(A,B)
// Writes j-th column of sum!4 with the new references:
sum!4[I,j] = C
// may eventually optimize since C is A
```

There are three kinds of arrays (for now all kept internally as one-dimensional arrays, but that's under discussion).

- Scalars: These are scalar variables we lift into arrays for the purposes of vectorization. For those, all writes are canonical writes and all reads are canonical reads. We may apply both raise dimension and drop dimension on these.
- Read-only input arrays: Read-only inputs. There are NO writes, while we may have non-canonical reads, $f(i, j, k)$. Phase 1 of Basic vectorization will add raise dimension operation at the beginning of the function and raise dimension may reshape arrays. If there are multiple "views" of the input array, there could be multiple raise dimension statements to create each one of these views. The invariant is that at reads in loops, the reads of "views" of the original input array will be canonical. Only raise dimension applies.
- Read-write output arrays: Writes are canonical (by restriction) but reads can be non-canonical. Conjecture: do nothing. Dependence analysis takes care of limiting vectorization so non-canonical access will work. Still work in progress. We may apply both raise and drop dimension.

4.3 Correctness Argument

4.4 Towards Extension of Basic Vectorization

4.4.1 Removal of Infeasible Edges

Array writes limit vectorization as they sometimes introduce infeasible loop-carried dependencies. Consider the following example: ANA: *Have to add citation to Aiken's paper*

    for i in range(N):
        A[i] = B[i] + 10;
        B[i] = A[i] * D[i-1];
        C[i] = A[i] * D[i-1];
        D[i] = B[i] * C[i];

In Cytron's SSA this code (roughly) translates into

    for i in range(N):
    1. $A_0 = \phi(A, A_1)$
    2. $B_0 = \phi(B, B_1)$
    3. $C_0 = \phi(C, C_1)$
    4. $D_0 = \phi(D, D_1)$
    5. $A_1 = A_0$; $A_1[i] = B_0[i] + 10$; {equiv. to Update}
    6. $B_1 = B_0$; $B_1[i] = A_1[i] * D_0[i-1]$;
    7. $C_1 = C_0$; $C_1[i] = A_1[i] * D_0[i-1]$;
    8. $D_1 = D_0$; $D_1[i] = B_1[i] * C_1[i]$;

There is a cycle around $B_0 = \phi(B, B_1)$ that includes statement $A_1[i] = B_0[i] + 10$; and that statement won't be vectorized even though in fact there is no loop-carried dependency from the write of $B_1[i]$ at 8 to the read of $... = B_0[i]$ at 6.

The following algorithm removes certain infeasible loop-carried dependencies that are due to array writes. Consider a loop with index $0 \le j < J$ nested at i, j, k. Here i represents the enclosing loops of j and k represents the enclosed loops in j.

> **for** each array A written in loop j **do**
>   { including enclosed loops in j }
>   dep = False
>   **for** each pair def: $A_m[f(i, j, k)] = ...$, and use: $... = A_n[f'(i, j, k)]$ in loop j **do**
>   **if** $\exists \underline{i}, j, \underline{j}', \underline{k}, \underline{k}'$, s.t. $0 \le \underline{i} < I$, $0 \le \underline{j}, \underline{j}' < J$, $0 \le \underline{k}, \underline{k}' < K$, $\underline{j} < \underline{j}'$, and $f(\underline{i}, \underline{j}, \underline{k}) = f'(\underline{i}, \underline{j}', \underline{k}')$ **then**
>   dep = True
>   **end if**
>   **end for**
>   **if** dep == False **then**
>   remove back edge into A's ϕ-node in loop j.
>   **end if**
> **end for**

ISHAQ: *Note to self: This algorithm is an instantiation for j loop, the one for k loop will be exactly the same, modulo variable name..*

Consider a loop j enclosed in some fixed $\underline{i}$. Only if an update (definition) $A_m[f(i, j, k)] = ...$ at some iteration $\underline{j}$ references the *same* array element as a use $... = A_n[f'(i, j, k)]$ at some later iteration $\underline{j}'$, we may have a loop-carried dependence for A due to this def-use pair. (In contrast, Cytron's

algorithm inserts a loop-carried dependency every time there is an array update.) The algorithm above examines all def-use pairs in loop j, including defs and uses in nested loops, searching for values $\underline{i}, \underline{j}, \underline{j}', \underline{k}, \underline{k}'$ that satisfy $f(\underline{i}, \underline{j}, \underline{k}) = f'(\underline{i}, \underline{j}', \underline{k}')$. If such values exist for some def-use pair, then there is a potential loop-carried dependence on A; otherwise there is not and we can remove the spurious backward edge thus "freeing up" statements for vectorization.

Consider the earlier example. There is a single loop, i. Clearly, there is no pair $\underline{i}$ and $\underline{i}'$, where $\underline{i} < \underline{i}'$ that make $\underline{i} = \underline{i}'$ (due to the def-use pairs of A 6-8 and 6-10). Therefore, we remove the back edge from 6 to 1. Analogously, we remove the back edges from 8 to 2 and 10 to 3. However, there are many values $\underline{i} < \underline{i}'$ that make $\underline{i} = \underline{i}' - 1$ and the back edge from 12 to 4 remains (def-use pairs for D). As a result of removing these spurious edges, Basic Vectorization will find that statement 6 is vectorizable. Statements 8, 10 and 12 will correctly appear in the FOR loop.

4.4.2 Array MUX refinement

ANA: *TODO: I think we should implement this.*

Next, the algorithm refines array MUX statements. MPC-source after Cytron's SSA may result in statements $A_j = MUX(..., A_k, A_l)$, which imply that any index of A can be written at this point and therefore there is a loop-carried dependency. In some cases the MUX can be refined to just a single index or a pair of indices, e.g., $A_j[i] = MUX(c, A_k[i], A_l[i])$.

This is to reduce the dimensionality of simd-ified computation. Technically, $A_j = MUX(..., A_k, A_l)$ is a simdified operation that can be carried out in parallel "in one round". However, particularly when A is a multi-dimensional array, there is substantial increase in the size of the arrays (vectors) we send to SIMD operations. Refining to an update to a specific index would reduce the size of those vectors. Note that this is a heuristic that handles a common case, but not all cases of array updates. ANA: *TODO: Simplify this algorithm, taking into account the restriction to canonical updates. It should be handling _all_ cases.*

> **for** each *stmt*: $A_j = MUX(c, A_k, A_l)$ in the MPC-source seq. **do**
> $i_1 = find\_update(A_k)$ { Is null when $A_k = \phi(...)$ }
> $i_2 = find\_update(A_l)$ { Is null when $A_l = \phi(...)$ }
>   **if** $i_1 == i_2$ or i_1 is null or i_2 is null **then**
>   { With our restrictions on writes we must have $i_1 = i_2$. }
>   replace *stmt* with
> $A_j = A_{j-1}$; $A_j[i_1] = MUX(c, A_k[i_1], A_l[i_1])$
>   **else**
>   *stmt* stays as is
>   **end if**
> **end for**

4.5 Extension of Basic Vectorization with Array Writes

4.5.1 Restricting Array Writes

For now, we restrict array updates to *canonical updates*. Assume (for simplicity) a two-dimensional array $A[I, J]$. A canonical update is the following:

```
for i in range(I):
    for j in range(J):
        ...
        A[i,j] = ...
        ...
```

The update $A[i, j]$ can be nested into an inner loop and there may be multiple updates, i.e., writes to $A[i, j]$. However, update such as $A[i - 1, j] = ...$ or $A[i - 1, j - 1] = ...$, etc., is not allowed. Additionally, while there could be several different loops that perform canonical updates, they must be of the same dimensionality, i.e., an update of higher or lower dimension, e.g., $A[i, j, k] = ...$ is not allowed. We compute the *canonical dimensionality* of each write array in the obvious way *before basic vectorization*. This restriction simplifies reasoning in this early stage of the compiler; we will look to relax the restriction in future work.

Another restriction/assumption is that we assume the output array is given as input with initial values, and it is of size consistent with its canonical dimensionality. ANA: *This restriction will require that we update our benchmarks that come with output arrays.*

Reads through an arbitrary formula, such as $A[i - 1]$ for example, are allowed; currently, our projection function returns dummy values if the read formula is out of bounds; we assume the programmer ensures that the program still computes correct output in this case.

4.5.2 Changes to Basic Vectorization

There are two changes to Basic Vectorization to account for write arrays. We can account for them during processing of raise dimension/drop dimension in Phase 1 of Basic Vectorization.

One change to Basic vectorization is the expansion of dimension if the array write or read occurs in a nested loop. That is, if there is an update $A[i, j] = ...$ that occurs in loop dimensionality i, j, k, $A[i, j]$ will be rewritten into $A[i, j, k]$. Similarly, a read $A[f(i, j)]$ will be rewritten into $A[f(i, j), k]$. ($A[f(i, j), k]$ will be reshaped as $f(i, j) * I * J + k$ during Phase 1 of Basic Vectorization. ANA: *TODO: Double check.*)

The other change concerns def-use edges $X \to Y$ where X defines and Y uses an array variable. The definition can be an update $A_2 = update(A_1, ...)$ or a pseudo ϕ-node $A_2 = \phi(A_0, A_1)$. Note that ϕ nodes for arrays are no subscript operations on those ϕ-nodes the way there are in scalar arrays. These edges are not handled in the same way as in Basic Vectorization, specifically, we do not raise and drop dimension as we do for scalars in Basic Vectorization. A key invariant is that the dimension of an array A cannot go above or drop below A's canonical dimensionality. We enumerate the cases of def-use edges.

(1) same-level $X \to Y$. We do nothing, just propagate the array, which happens to be of the right dimension. ANA: *There might be some opportunities to do copy propagation optimization and save some cycles, but let's leave this for later.*

(2) inner-to-outer $X \to Y$. If dimensionality of the loop enclosure of X is greater than the canonical dimensionality of the array, then add *drop_dim(...)* at Y, as in Basic Vectorization. Otherwise, do nothing.

(3) outer-to-inner $X \to Y$. If dimensionality of loop enclosure of Y is greater than the canonical dimensionality of the array, then add *raise_dim(...)* (at X) as in Basic Vectorization. Otherwise, do nothing.

(4) "mixed" $X \to Y$. We assume that the mixed edge is transformed into an inner-to-outer followed by outer-to-inner edge before we perform vectorization, just as with Basic vectorization.

4.5.3 Examples with Array Writes

Example 1. First, the canonical dimensionality of all A, B, C and D is 1. Thus, there is no addition of extra dimension for inner loops. After Phase 1 of Basic Vectorization the Aiken's array write example will be (roughly) as follows:

for i in range(N):
1. $A_1 = \phi(A_0, A_2)$
2. $B_1 = \phi(B_0, B_2)$
3. $C_1 = \phi(C_0, C_2)$
4. $D_1 = \phi(D_0, D_2)$
5. $A_2 = update(A_1, \vec{i}, B_1[\vec{i}] + 10);$
6. $B_2 = update(B_1, \vec{i}, A_2[\vec{i}] * D_1[i - 1]);$
7. $C_2 = update(C_1, \vec{i}, A_2[\vec{i}] * D_1[i - 1]);$
8. $D_2 = update(D_1; \vec{i}, B_2[\vec{i}] * C_2[\vec{i}]);$

Phase 2 computes the closure of 4; $cl = \{4, 6, 7, 8\}$ while 5 is vectorizable. Recall that 1,2, and 3 are targetless phi-nodes. Since the closure cl includes updates to B and C, the corresponding phi-nodes are added to the closure, i.e., the FOR loop. The uses of A_1 and B_1 in the vectorized statement turn into uses of A_0 and B_0 respectively. (But note that A_0 is irrelevant; the updates writes into array A_2 in parallel.)

1. $A_2 = update(A_0; \vec{i}, ADD\_SIMD(B_0[\vec{i}], [10, ...])$ {Fully vectorized, size N.}
FOR i=0; i¡N; i++; { MOTION loop }
2. $B_1 = \phi(B_0, B_2)$
3. $C_1 = \phi(C_0, C_2)$
4. $D_1 = \phi(D_0, D_2)$
5. $B_2 = update(B_1, i, MUL(A_2[i], D_1[i - 1])$
6. $C_2 = update(C_1, i, MUL(A_2[i], D_1[i - 1])$
7. $D_2 = update(D_1, i, MUL(B_2[i], C_1[i])$

ANA: *TODO: Format either Histogram or Matrix Multiplication. Both are interesting examples that use output arrays.*

5 DIVIDE-AND-CONQUER

ANA: *TODO: Now that we have broken FOR loops into smaller chunks, we can add Divide-and-conquer reasoning with Z3 and implement this additional transform.*

6   IMPLEMENTATION AND EVALUATION

7   FUTURE WORK

8   CONCLUSIONS